

Final Exam

Pranay Kumar Kodeboyina

2023-05-07

```
# importing the Necessary libraries  
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
## filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(leaps)  
library(dbscan)
```

```
##  
## Attaching package: 'dbscan'
```

```
## The following object is masked from 'package:stats':  
##  
## as.dendrogram
```

```
library(esquisse)  
library(readr)
```

```
#Extracting the current working directory
getwd()
```

```
## [1] "/Users/kodeboyina/Documents/Kent State/Sem1/Fundamentals of ML/Final Exam"
```

```
#Loading PUDL csv data Import the data set into R
Fuel <- read.csv("data/fuel_receipts_costs.csv", header = TRUE, sep = ",", stringsAsFactors = TRUE)

#Observing the first 10 observations of the data set
head(Fuel, n=10L)
```

ro...	plant_id_eia	report_date	contract_type_code	contract_expiration_date	ener	
<int>	<int>	<fct>	<fct>	<fct>	<fct>	
1	1	3	2008-01-01	C	2008-04-01	BIT
2	2	3	2008-01-01	C	2008-04-01	BIT
3	3	3	2008-01-01	C		NG
4	4	7	2008-01-01	C	2015-12-01	BIT
5	5	7	2008-01-01	S	2008-11-01	BIT
6	6	7	2008-01-01	S	2008-01-01	BIT
7	7	7	2008-01-01	S		NG
8	8	8	2008-01-01	C	2008-12-01	BIT
9	9	8	2008-01-01	C	2008-03-01	BIT
10	10	8	2008-01-01	C	2010-12-01	BIT

1-10 of 10 rows | 1-7 of 24 columns

```
#Observing the data sets that have missing values more than 50 percent and omitting those columns in the analysis
```

```
# calculate the percentage of missing values in each column
missing_pct <- colMeans(is.na(Fuel)) * 100
```

```
# create a data frame with the missing percentages as a single column
missing_df <- data.frame(percent_missing = missing_pct)
```

```
# print the resulting data frame
print(missing_df)
```

```
##                                percent_missing
## rowid                        0.00000
## plant_id_eia                 0.00000
## report_date                  0.00000
## contract_type_code           0.00000
## contract_expiration_date     0.00000
## energy_source_code           0.00000
## fuel_type_code_pudl          0.00000
## fuel_group_code              0.00000
## mine_id_pudl                 64.40506
## supplier_name                0.00000
## fuel_received_units          0.00000
## fuel_mmbtu_per_unit          0.00000
## sulfur_content_pct           0.00000
## ash_content_pct              0.00000
## mercury_content_ppm          47.56805
## fuel_cost_per_mmbtu          32.90369
## primary_transportation_mode_code 0.00000
## secondary_transportation_mode_code 0.00000
## natural_gas_transport_code    0.00000
## natural_gas_delivery_contract_type_code 0.00000
## moisture_content_pct         84.88639
## chlorine_content_ppm         84.88639
## data_maturity                0.00000
```

#Upon observing the data we can see that the following columns has missing values more than 50% - mine_id_pudl, moisture_content_pct, chlorine_content_ppm

```
# set the random seed for reproducibility
set.seed(8627)
```

```
# create a new data set without the specified columns and considering the columns without the specified columns
Fuel_data_ALL <- Fuel[, -c(1,2,3,5,9,10,15,18,19,20,21,22,23)]
head(Fuel_data_ALL)
```

contract_type_code <fct>	energy_source_code <fct>	fuel_type_code_pudl <fct>	fuel_group_code <fct>	
1 C	BIT	coal	coal	1
2 C	BIT	coal	coal	
3 C	NG	gas	natural_gas	
4 C	BIT	coal	coal	
5 S	BIT	coal	coal	
6 S	BIT	coal	coal	

6 rows | 1-6 of 11 columns

```
#head(Fuel_data_ALL)

# calculate the number of rows to select (2% of the total rows)
nrows <- round(nrow(Fuel_data_ALL) * 0.02)

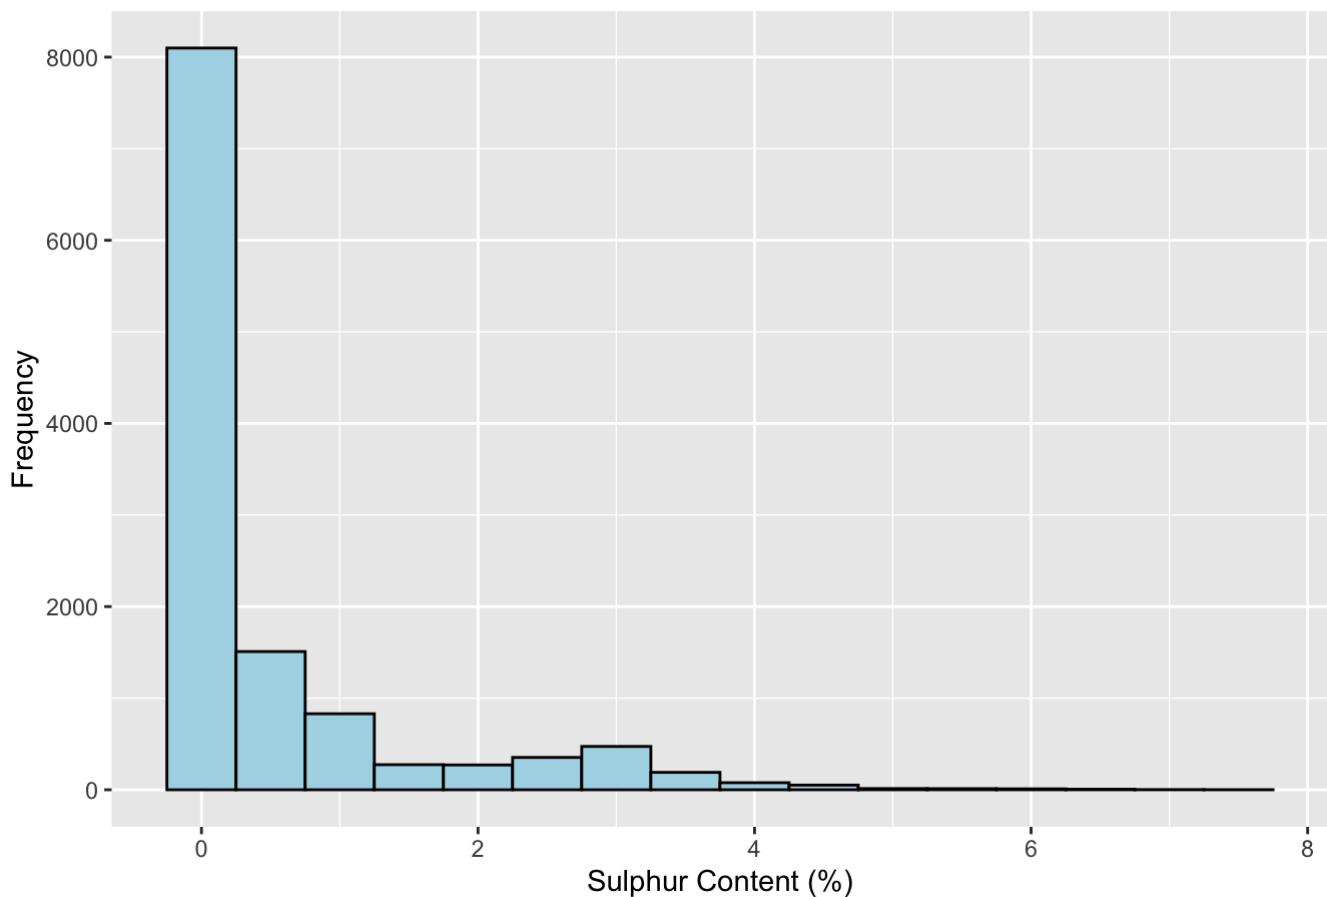
# randomly select the rows
idx <- sample(nrow(Fuel_data_ALL), nrows, replace = FALSE)

# create the new data set with the randomly selected rows
Fuel_data <- Fuel_data_ALL[idx, ]

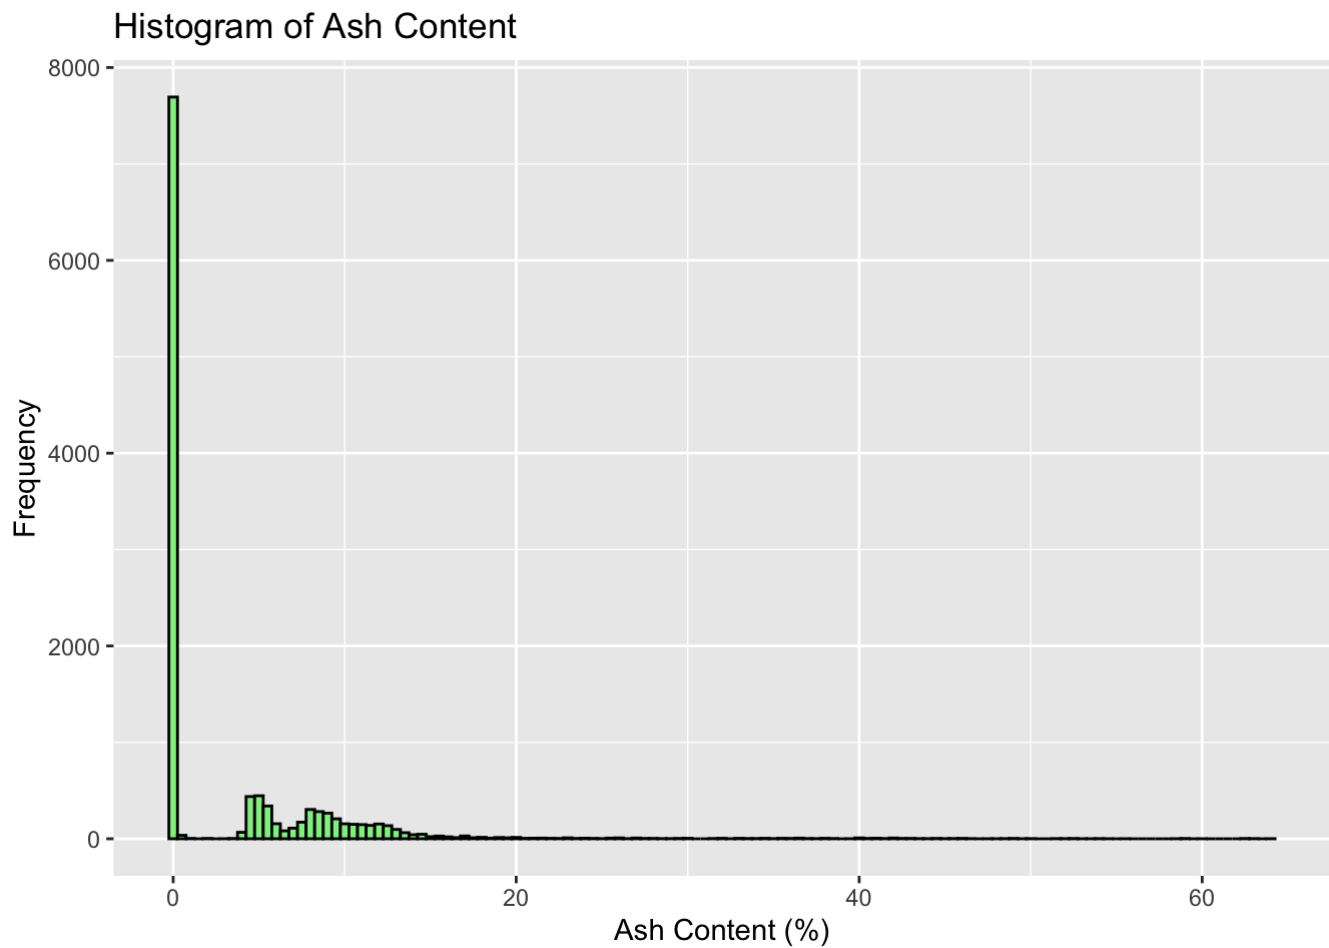
#View(Fuel_data)

#Exploratory data Analysis
# Create histograms of sulphur and ash content
ggplot(Fuel_data, aes(x=sulfur_content_pct)) +
  geom_histogram(binwidth=0.5, color="black", fill="lightblue") +
  labs(x="Sulphur Content (%)", y="Frequency", title="Histogram of Sulphur Content")
```

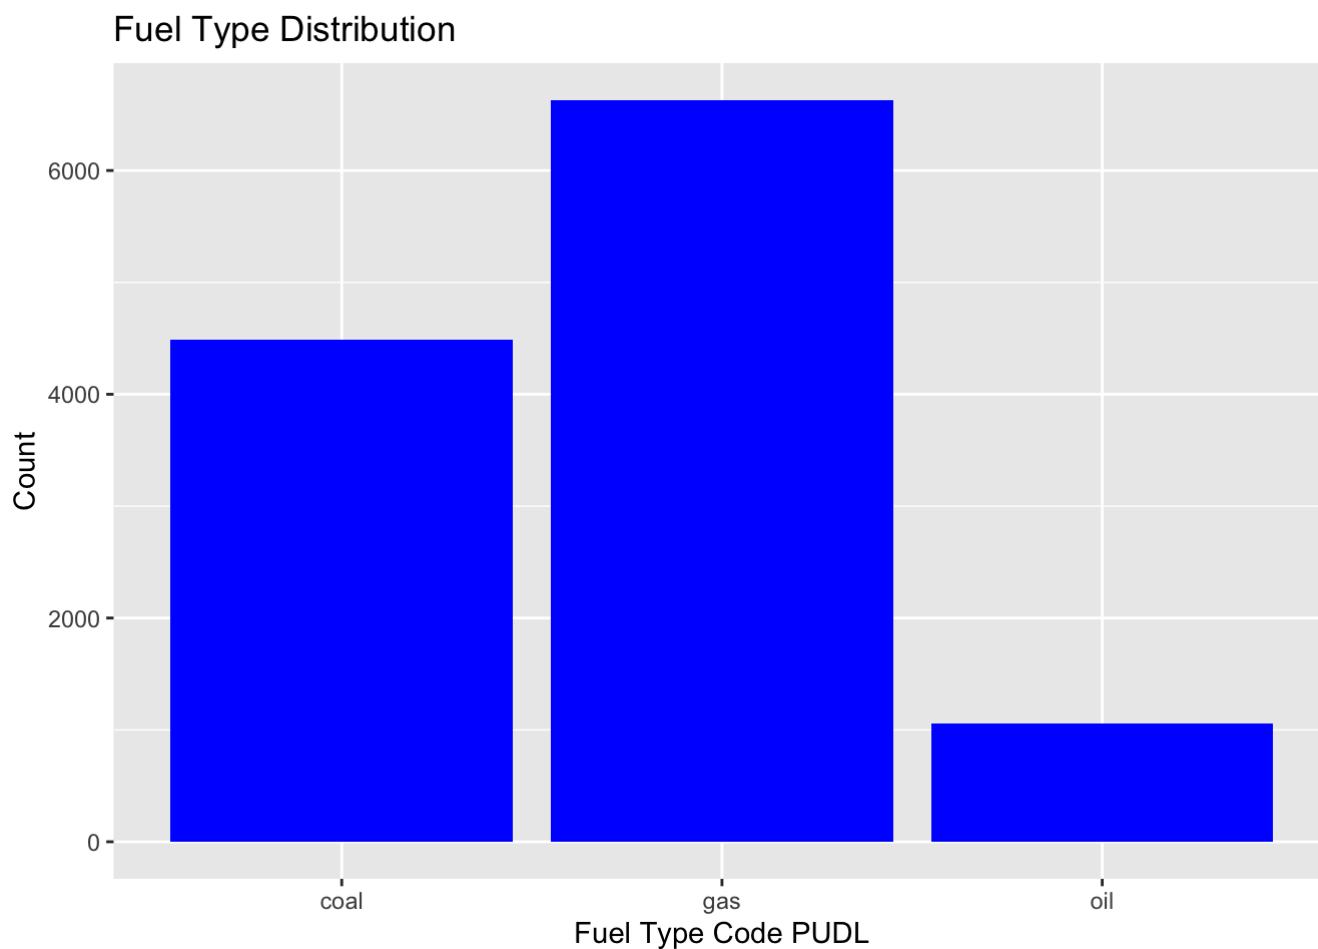
Histogram of Sulphur Content



```
ggplot(Fuel_data, aes(x=ash_content_pct)) +
  geom_histogram(binwidth=0.5, color="black", fill="lightgreen") +
  labs(x="Ash Content (%)", y="Frequency", title="Histogram of Ash Content")
```



```
# create a bar graph of the fuel_type_code_pudl categories
ggplot(Fuel_data, aes(x = fuel_type_code_pudl)) +
  geom_bar(fill = "blue") +
  labs(title = "Fuel Type Distribution",
       x = "Fuel Type Code PUDL",
       y = "Count")
```



```
# Create dummy variables for fuel_type_code_pudl to create categorical in to numeric variables
Fuel_data <- Fuel_data %>%
  mutate(fuel_type_gas = ifelse(fuel_type_code_pudl == "gas", 1, 0),
         fuel_type_coal = ifelse(fuel_type_code_pudl == "coal", 1, 0),
         fuel_type_oil = ifelse(fuel_type_code_pudl == "oil", 1, 0))

# Remove the original fuel_type_code_pudl column
Fuel_data <- select(Fuel_data, -fuel_type_code_pudl)

# We can find the missing values in fuel_cost_per_mmbtu data and replacing them with the median values for the analysis

# Calculate the median value of fuel_cost_per_mmbtu
median_value <- median(Fuel_data$fuel_cost_per_mmbtu, na.rm = TRUE)

# Replace missing values with median
Fuel_data$fuel_cost_per_mmbtu[is.na(Fuel_data$fuel_cost_per_mmbtu)] <- median_value
```

```
# Generate a vector of random indices for the data frame and splitting 75% of data
train_idx <- sample(nrow(Fuel_data), round(nrow(Fuel_data) * 0.75), replace = FALSE)

# Create the training set
train_data <- Fuel_data[train_idx, ]

# Create the test set by excluding the training set indices
test_data <- Fuel_data[-train_idx, ]

nrow(train_data)
```

```
## [1] 9128
```

```
nrow(test_data)
```

```
## [1] 3043
```

```
library(caret)

# Create a preprocessing object using the train_data dataset
preproc_obj <- preProcess(train_data[, 4:7], method = c("center", "scale"))

# Use the preprocessing object to normalize the train_data and test_data datasets
train_data_norm <- predict(preproc_obj, train_data)
test_data_norm <- predict(preproc_obj, test_data)

#View(train_data_norm)
```

#We have divided the data in to Train and Test set and we are finding Optimal values of K using Gap stat method and Silhouette method

#In order determine the optimal value of k, employing different methods to determine k

```
library(cluster)
```

```
library(ggplot2)
```

```
wss <- c()
```

```
for (i in 1:10) {
```

```
  kmeans_fit <- kmeans(train_data_norm[, 4:7], i, nstart = 25)
```

```
  wss[i] <- kmeans_fit$tot.withinss
```

```
}
```

```
elbow_df <- data.frame(K = 1:10, WSS = wss)
```

```
ggplot(elbow_df, aes(x = K, y = WSS)) +
```

```
  geom_line() +
```

```
  geom_point() +
```

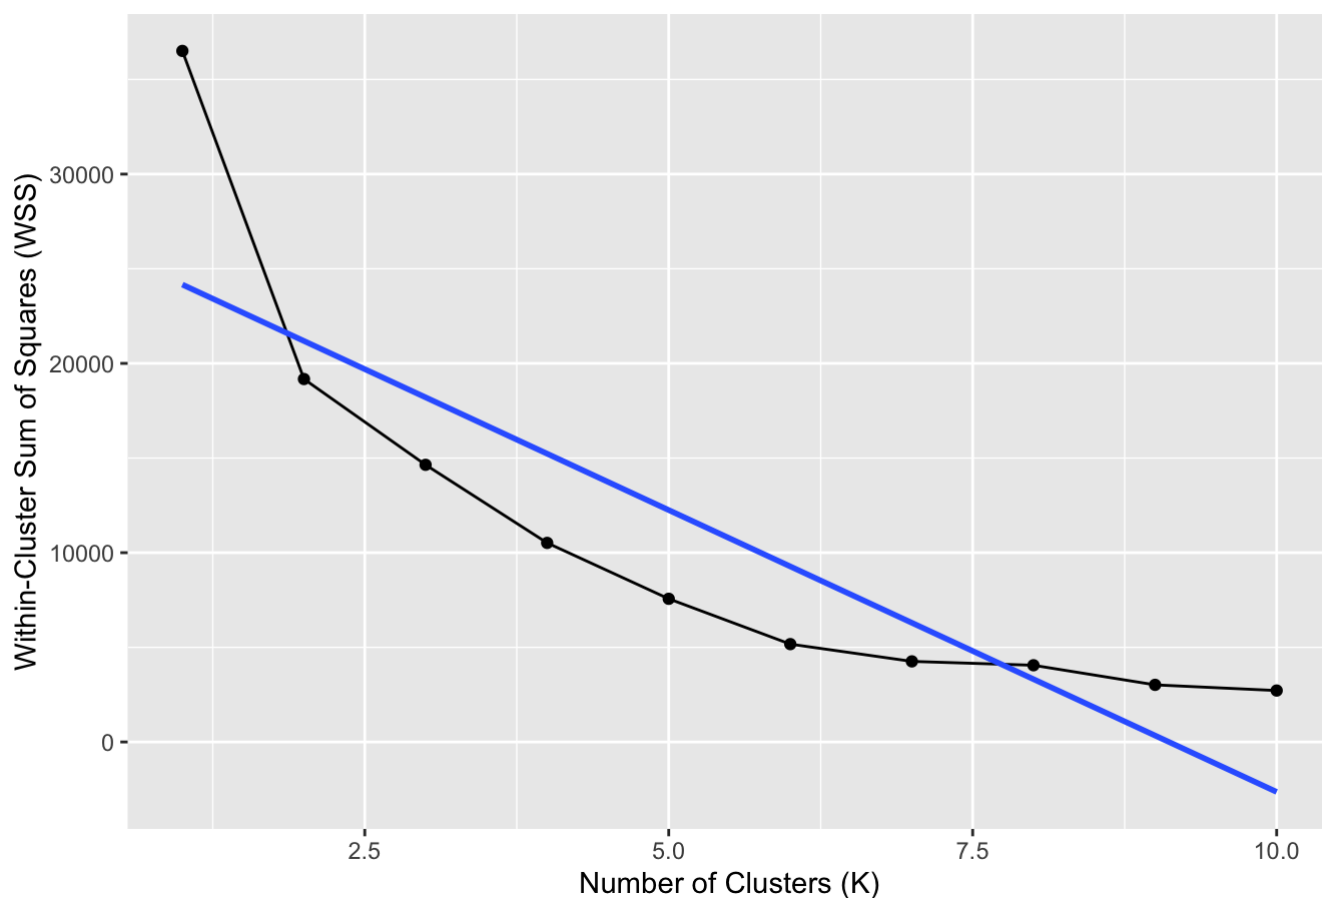
```
  labs(x = "Number of Clusters (K)", y = "Within-Cluster Sum of Squares (WSS)") +
```

```
  ggtitle("Elbow Method with Slope") +
```

```
  geom_smooth(method = "lm", se = FALSE)
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

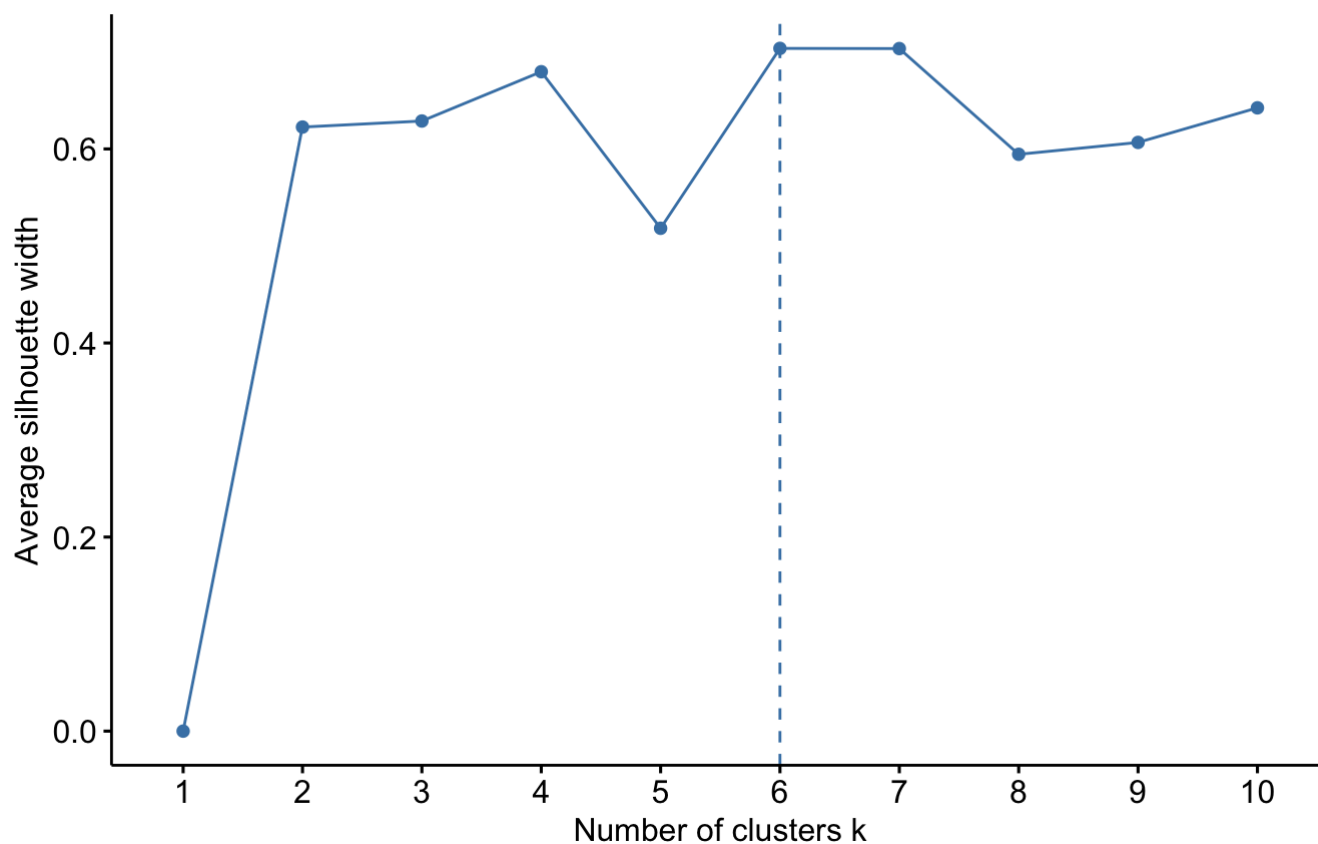
Elbow Method with Slope




```
fviz_nbclust(train_data_norm[, 4:7], kmeans, method = "silhouette") + labs(subtitle =
"Silhouette Method To Determine Optimal Value of K")
```

Optimal number of clusters

Silhouette Method To Determine Optimal Value of K



#From the above method we can see that using Elbow method the value of K is 5 and the value of K is 5 with the silhouette method. I am choosing K value of 5 to cluster the data as it has produced clusters with clear gap between each other.

```
kmean <- kmeans(train_data_norm[, 4:7], centers = 6)
```

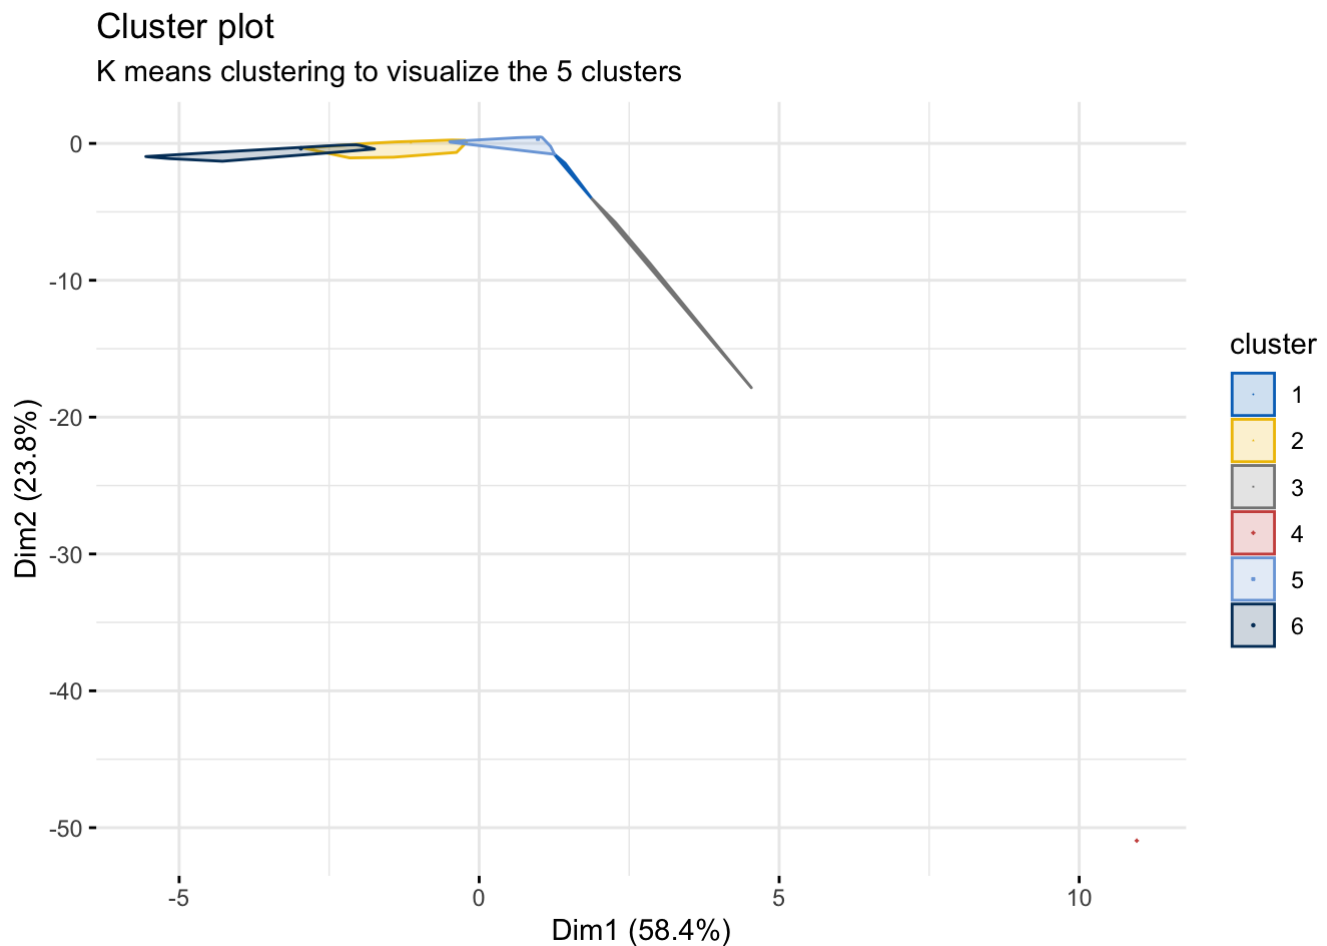
```
## Below finding cluster center for all rows and columns
kmean$centers
```

```
##    fuel_received_units fuel_mmbtu_per_unit sulfur_content_pct ash_content_pct
## 1          2.1677469         -0.7963608         -0.51998098         -0.5500416
## 2         -0.2211667          1.1873022          0.08285029          0.6310875
## 3          6.4890043         -0.7996130         -0.51998098         -0.5500416
## 4         52.1791803         -0.8911217         -0.51998098         -0.5500416
## 5         -0.1472000         -0.7231877         -0.49352278         -0.5500416
## 6         -0.2424190          1.3939162          2.25568747          1.5256845
```

```
## Number of observation in each cluster
kmean$size
```

```
## [1] 498 2156 62 1 5219 1192
```

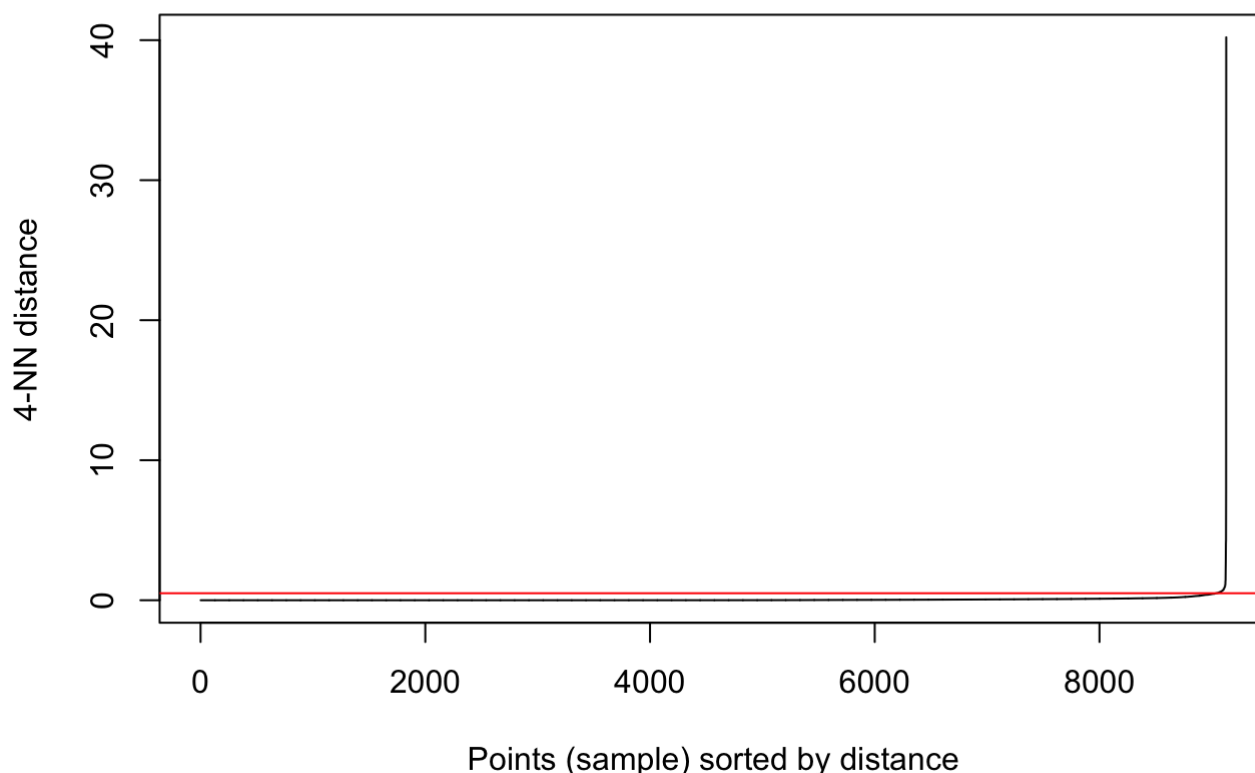
```
library(factoextra)
##For the above observations applying the K means clustering to visualize the 5 clusters
fviz_cluster(kmean, data=train_data_norm[,4:7], geom = "convex",
              palette = "jco", ggtheme = theme_minimal()) + labs(subtitle = "K means clustering to visualize the 5 clusters")
```



After using K mean algorithm we can see that clusters 3,4 and 5 are overlapping hence it indicates that the data is not well-suited for this type of clustering. In such cases, we can consider using density-based clustering algorithms such as DBSCAN. Compared to k-means clustering, DBSCAN is more robust to noise and can handle data with varying densities and shapes

```
library(dbSCAN)

dbSCAN::kNNdistplot(train_data_norm[,4:7], k = 4)
abline(h = 0.5,col="red")
```



#The kNN distplot with k=4 has identified a clear elbow point at 0.5, indicating a reasonable value for the epsilon parameter in the DBSCAN algorithm. This means that data points that are within a distance of 0.5 from each other are considered to be in the same cluster. The appropriate value for minPts is then chosen based on domain knowledge and experimentation.

```
# Perform DBSCAN clustering
```

```
#Selecting numerical data to form clusters:
```

```
Training_numerical<-train_data[,c(4:7)]
```

```
#Normalizing the data:
```

```
Training_norm<-scale(Training_numerical)
```

```
db <- dbscan::dbscan(Training_norm, eps = 0.5, minPts = 25)
```

```
db
```

```
## DBSCAN clustering for 9128 objects.
```

```
## Parameters: eps = 0.5, minPts = 25
```

```
## Using euclidean distances and borderpoints = TRUE
```

```
## The clustering contains 3 cluster(s) and 294 noise points.
```

```
##
```

```
##      0      1      2      3
```

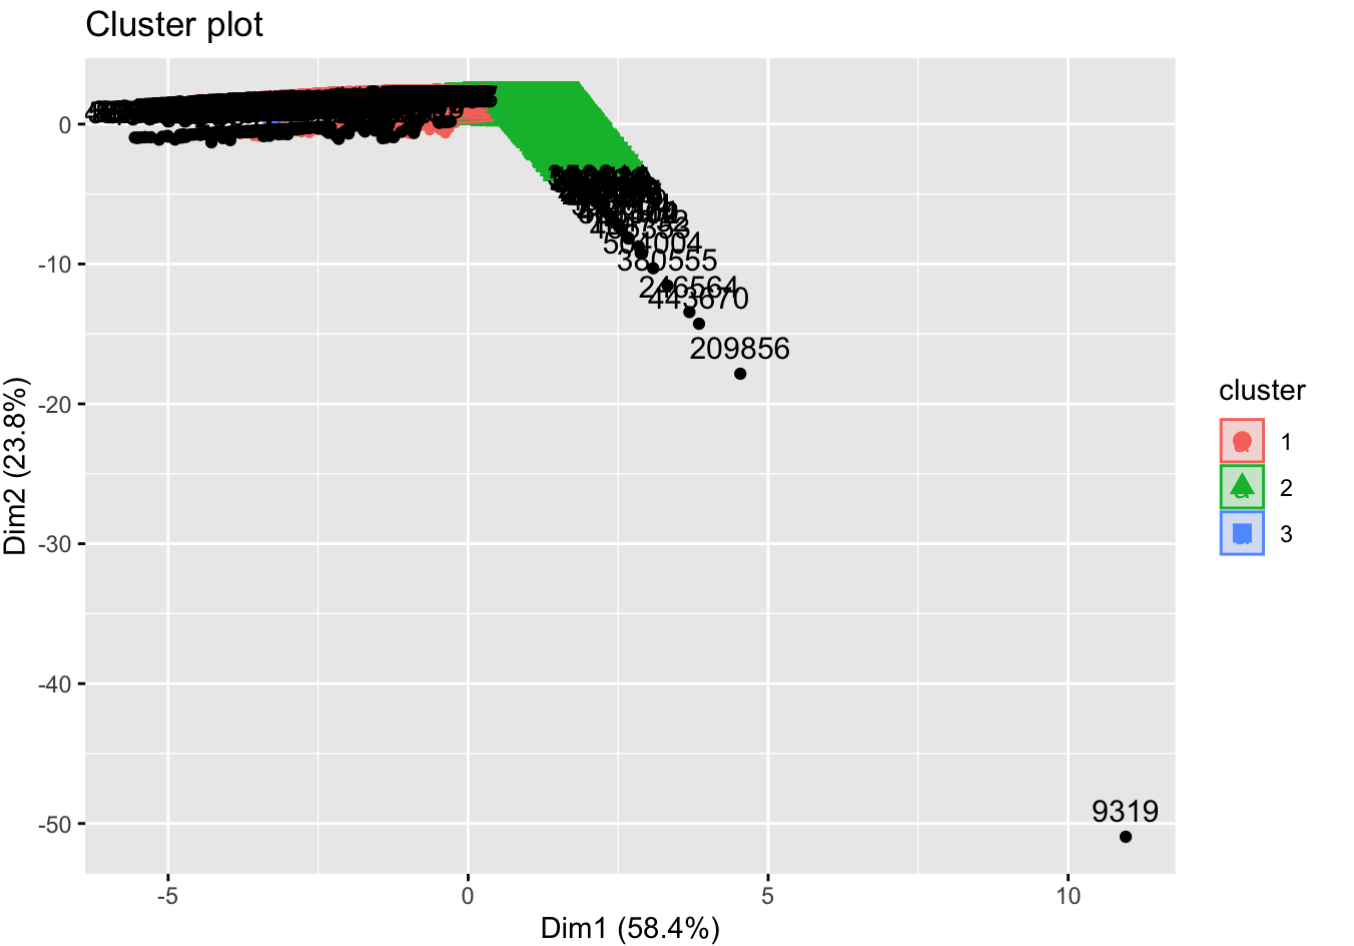
```
## 294 3072 5737    25
```

```
##
```

```
## Available fields: cluster, eps, minPts, dist, borderPoints
```

```
#The result of the clustering is that there are 4 clusters and 256 noise points. The clusters are labeled as 0, 1, 2, and 3, with cluster 0 containing 294 points, cluster 1 containing 3072 points, cluster 2 containing 5737 points, and clusters 3 has 30 points.

# Visualize the clustering results
library(factoextra)
fviz_cluster(db, train_data_norm[,4:7], stand = TRUE, show.clust.cent = TRUE, ellipse.type = "norm")
```



```
#Let us explore the clusters formed and try to understand how each attribute is behaving in different cluster.

#Assigning clusters to the original data:
train_set<-cbind(train_data,db$cluster)

head(train_set)
```

	contract_type_code <fct>	energy_source_code <fct>	fuel_group_code <fct>	fuel_received_units <dbl>
301641	C	BIT	coal	7603
84826	S	NG	natural_gas	85
581296	S	NG	natural_gas	5911

contract_type_code <fct>	energy_source_code <fct>	fuel_group_code <fct>	fuel_received_units <dbl>
438751 S	RFO	petroleum	2359
258708 C	WC	coal	17840
592428 C	NG	natural_gas	283417

6 rows | 1-5 of 14 columns

#Let us explore the clusters formed and try to understand how each attribute is behaving in different cluster.

```
library(dplyr)
#Finding mean within each cluster to interpret the clusters:
train_set%>%group_by(db$cluster)%>%
  summarize(avg_units=mean(fuel_received_units),
            avg_cost=mean(fuel_cost_per_mmbtu),
            avg_mmbtu=mean(fuel_mmbtu_per_unit))
```

db\$cluster <int>	avg_units <dbl>	avg_cost <dbl>	avg_mmbtu <dbl>
0	1008327.02	3.031177	15.285986
1	45303.13	2.691056	21.660984
2	333651.24	7.235618	1.672761
3	7182.60	3.232680	12.457120

4 rows

#From the above observation we can see that Average cost for cluster 2 is the highest and maximum Average heat is generated by cluster 1 fuel type and least heat is generated by Cluster 2 unit fuel

#Adding the cluster information to original data without dummy variable and let us use this for further analysis.

```
set.seed(8627)
Part_data<-caret::createDataPartition(Fuel_data$fuel_mmbtu_per_unit,p=0.75,list = FALSE)

Data_final<-Fuel_data[train_idx,]

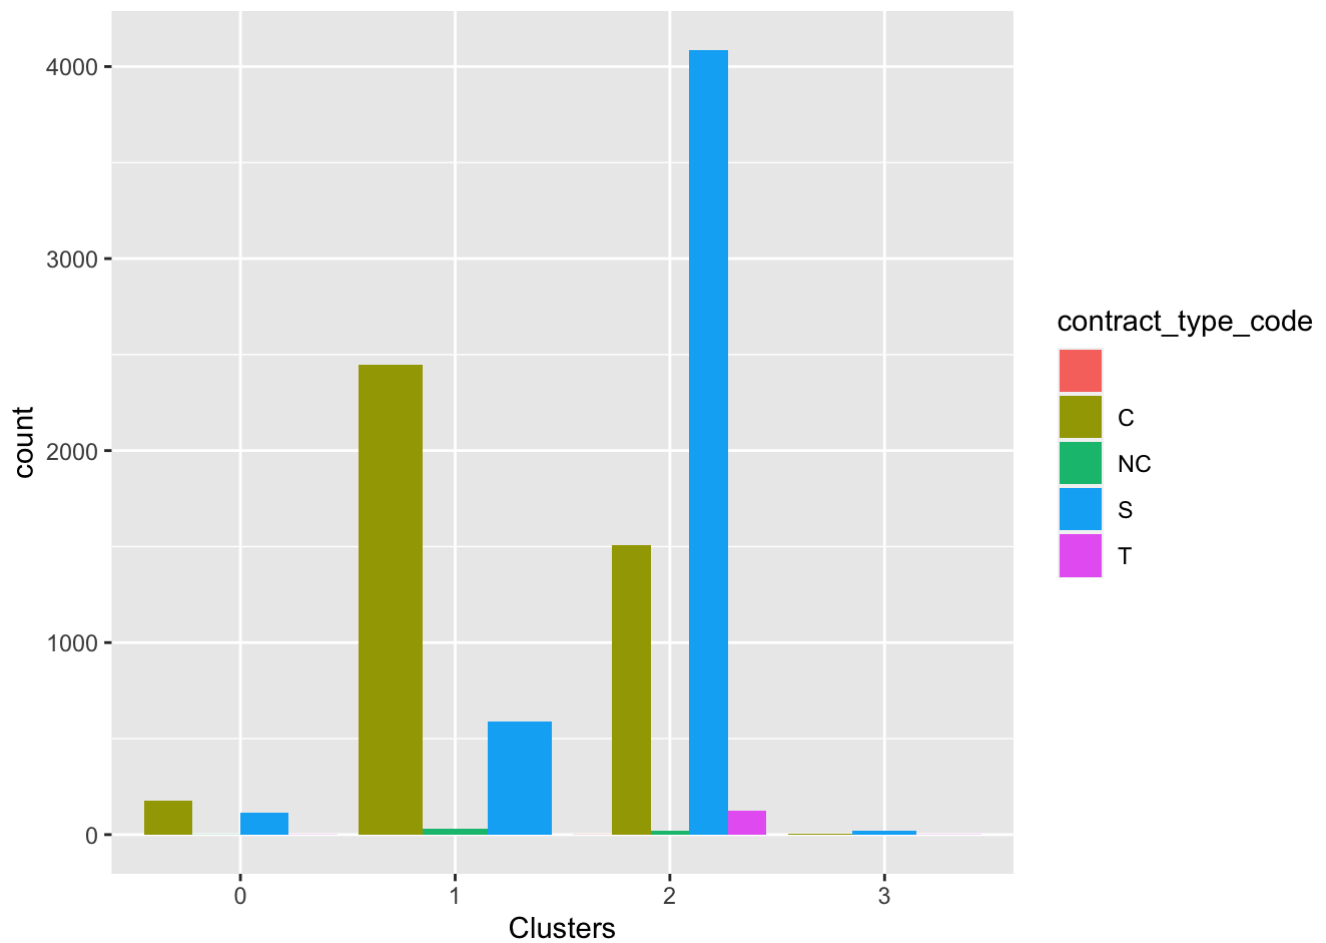
nrow(Data_final)
```

```
## [1] 9128
```

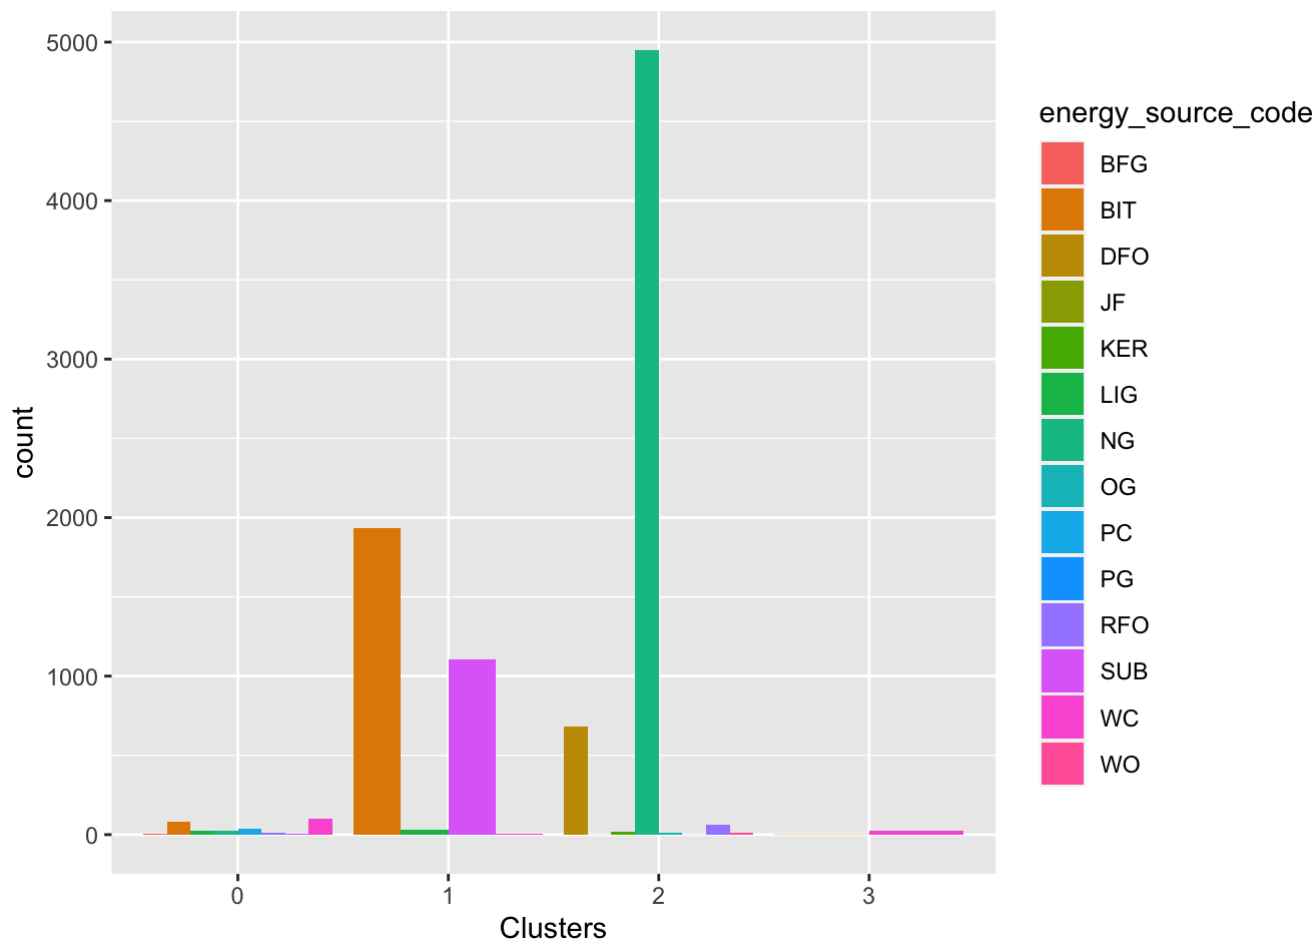
```
# Cluster 1 - Gas  
# Cluster 2 - Oil  
# Cluster 3 - Coal  
  
Data_final$cluster<-db$cluster
```

```
library(ggplot2)
```

```
#Observing the Distribution of Fuel across Different Contracts  
ggplot(Data_final, mapping = aes(factor(cluster), fill =contract_type_code))+geom_bar  
(position='dodge')+labs(x = 'Clusters')
```

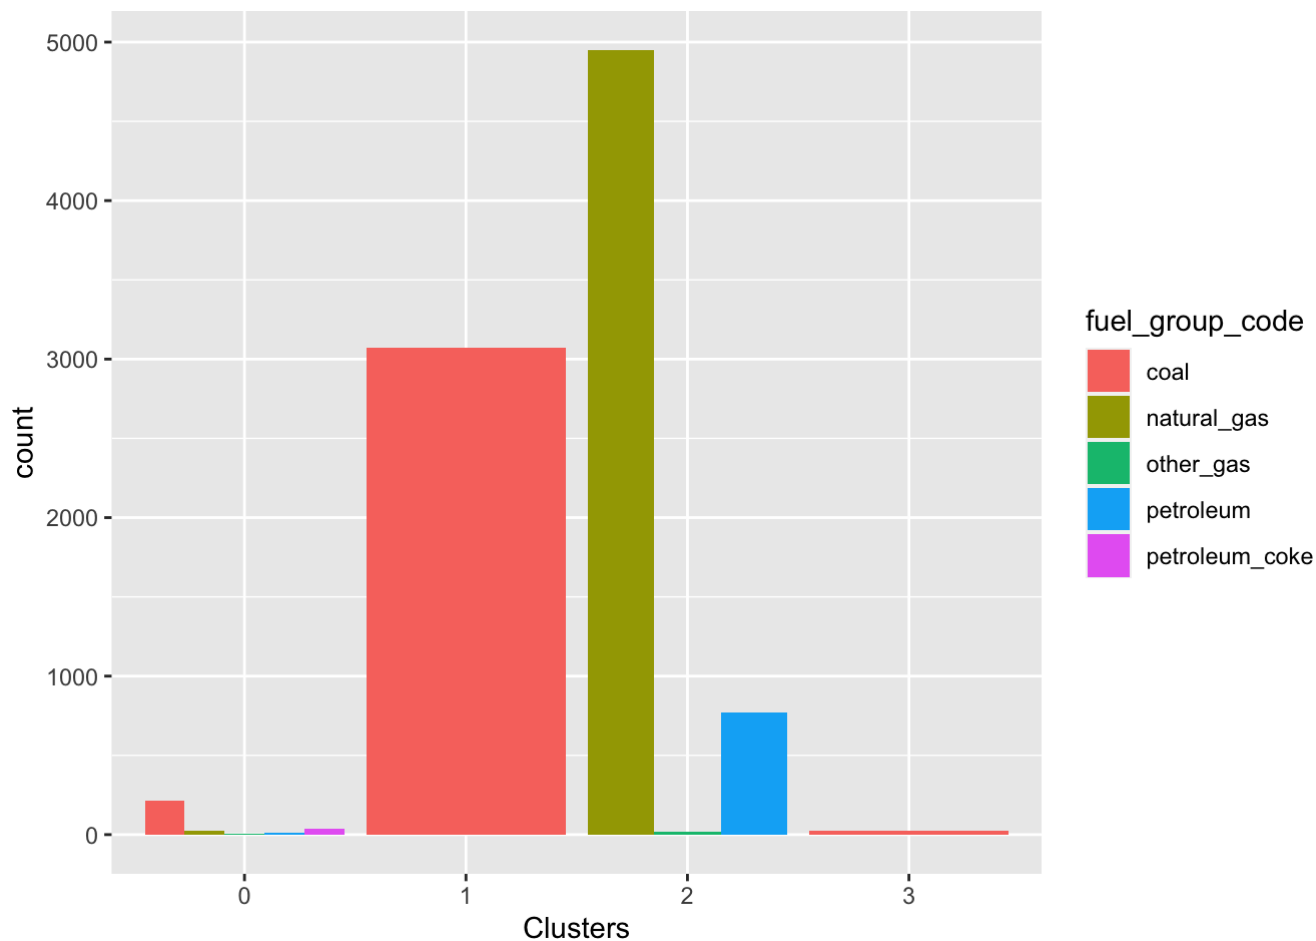


```
#Observing the Distribution of Fuel across Different Energy sources  
ggplot(Data_final, mapping = aes(factor(cluster), fill =energy_source_code))+geom_bar  
(position='dodge')+labs(x = 'Clusters')
```

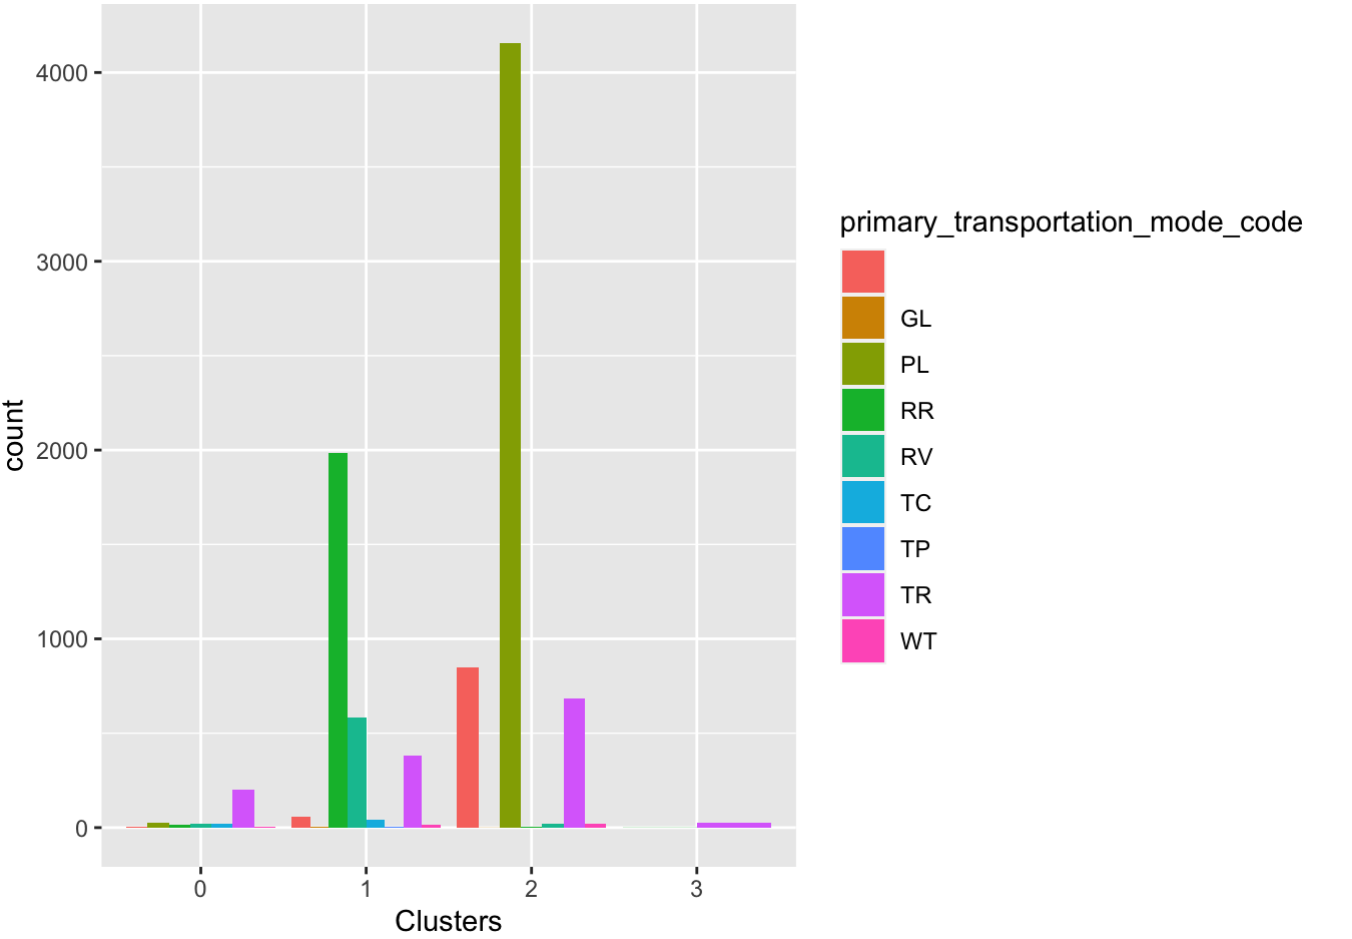


#Observing the Distribution of Fuel across Different Fuel group codes

```
ggplot(Data_final, mapping = aes(factor(cluster), fill =fuel_group_code))+geom_bar(position='dodge')+labs(x = 'Clusters')
```



```
#Observing the Distribution of Fuel across Different Transportation codes  
ggplot(Data_final, mapping = aes(factor(cluster), fill =primary_transportation_mode_c  
ode))+geom_bar(position='dodge')+labs(x ='Clusters')
```

Cluster 1 - Gas

#The analysis provides insights into the characteristics of Cluster 1, which represents the Gas fuel type.

#The analysis suggests that Gas has the lowest average fuel cost per mmbtu compared to other fuel types, which could make it a cost-effective option for energy generation.

#Gas has the lowest average fuel mmbtu per unit, indicating that it generates less heat content per unit compared to other fuel types.

#Gas does not contain any ash, sulfur, and mercury content, which could make it a clean and environmentally friendly fuel option.

#Natural gas is the energy source code for Gas, and pipelines are the most commonly used transportation type to supply this fuel type. This information could be useful in assessing the infrastructure requirements and associated costs for using Gas as a fuel type.

#The graph indicates that Gas is mainly purchased on spot rather than through contracts, which could be an important factor to consider for procurement strategies.

Cluster 2 - Oil

#Oil is the most expensive type of fuel and is purchased only on spot. This suggests that the buyers are not willing to enter into long-term contracts for purchasing oil, possibly due to the high cost and uncertainty of future prices.

#Oil has a relatively low average units received in comparison to gas and coal could be due to its high cost. Customers may prefer to use other fuel types that are cheaper and more cost-effective.

#Sulfur content in oil is important as it indicates that oil may have negative environmental impacts such as air pollution

Cluster 3 - Coal

#Coal has the lowest average cost per unit of fuel, it also has the highest average fuel mmbtu per unit, indicating that it produces more heat energy per unit than gas or oil. This could explain why it is widely used in the US despite its environmental impact.

#Coal is purchased both on spot and on contract, but there are more spot purchases than contract purchases.

#Coal include BIT and SUB, which stand for Bituminous Coal and Sub bituminous Coal, respectively. These are two types of coal that differ in their energy content and chemical properties, and they are both widely used in the US for electricity generation and other industrial purposes.

#Presence of ash, sulfur, and mercury in coal can have significant environmental and health impacts, including air pollution, acid rain, and water pollution. It is important to note that the use of coal is declining in the US and many other countries due to these impacts, as well as the increasing availability and affordability of renewable energy sources.

#Use multiple-linear regression to determine the best set of variables to predict fuel_cost_per_mmbtu. Running the multiple linear regression model to determine the best set of variables to predict fuel_cost_per_mmbtu by considering variables which were used to form clusters:

```
Model_data<- lm(train_data$fuel_cost_per_mmbtu~.,data=train_data)
```

```
summary(Model_data)
```

```
##
## Call:
## lm(formula = train_data$fuel_cost_per_mmbtu ~ ., data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.1    -3.5    -1.1     0.2   4640.2
##
## Coefficients: (7 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.763e+01  6.338e+01   0.278   0.781
## contract_type_codeC    3.120e+00  5.348e+01   0.058   0.953
## contract_type_codeNC    3.462e+00  5.398e+01   0.064   0.949
## contract_type_codeS    4.208e+00  5.347e+01   0.079   0.937
## contract_type_codeT    6.608e-01  5.368e+01   0.012   0.990
## energy_source_codeBIT   -2.051e+01  3.979e+01  -0.516   0.606
## energy_source_codeDFO   -9.914e+00  3.486e+01  -0.284   0.776
## energy_source_codeJF    -2.109e+01  6.352e+01  -0.332   0.740
## energy_source_codeKER   -1.898e+01  3.691e+01  -0.514   0.607
## energy_source_codeLIG   -1.987e+01  3.732e+01  -0.532   0.595
## energy_source_codeNG    -1.635e+01  3.395e+01  -0.481   0.630
## energy_source_codeOG    -1.849e+01  3.683e+01  -0.502   0.616
## energy_source_codePC    -2.086e+01  4.177e+01  -0.499   0.617
## energy_source_codePG    -8.181e+00  6.368e+01  -0.128   0.898
## energy_source_codeRFO   -1.590e+01  3.519e+01  -0.452   0.651
## energy_source_codeSUB   -2.073e+01  3.715e+01  -0.558   0.577
## energy_source_codeWC    -2.052e+01  3.949e+01  -0.520   0.603
## energy_source_codeWO    -1.804e+01  3.812e+01  -0.473   0.636
## fuel_group_codenatural_gas      NA          NA      NA      NA
## fuel_group_codeother_gas      NA          NA      NA      NA
## fuel_group_codepetroleum      NA          NA      NA      NA
## fuel_group_codepetroleum_coke    NA          NA      NA      NA
## fuel_received_units    -9.372e-07  6.936e-07  -1.351   0.177
## fuel_mmbtu_per_unit     9.255e-02  7.333e-01   0.126   0.900
## sulfur_content_pct     -1.181e-01  1.101e+00  -0.107   0.915
## ash_content_pct        2.618e-02  2.874e-01   0.091   0.927
## primary_transportation_mode_codeGL -1.922e+00  2.461e+01  -0.078   0.938
## primary_transportation_mode_codePL  2.085e+00  1.993e+00   1.046   0.296
## primary_transportation_mode_codeRR  3.967e-01  5.858e+00   0.068   0.946
## primary_transportation_mode_codeRV -6.066e-02  6.068e+00  -0.010   0.992
## primary_transportation_mode_codeTC  7.305e-01  9.135e+00   0.080   0.936
## primary_transportation_mode_codeTP  7.199e-01  2.461e+01   0.029   0.977
## primary_transportation_mode_codeTR  2.134e-01  5.861e+00   0.036   0.971
## primary_transportation_mode_codeWT  1.682e+00  1.001e+01   0.168   0.867
## fuel_type_gas          NA          NA      NA      NA
## fuel_type_coal         NA          NA      NA      NA
## fuel_type_oil          NA          NA      NA      NA
##
## Residual standard error: 53.46 on 9098 degrees of freedom
## Multiple R-squared:  0.002928, Adjusted R-squared:  -0.0002497
## F-statistic: 0.9214 on 29 and 9098 DF, p-value: 0.5867
```

#Fuel received units, fuel_type_coal and fuel_type_oil best determine the fuel_cost_per_mmbtu variable.

#Checking the prediction of the above model on Test data

#We could see that the difference between predicted values and actual values is too high, which means that it is too complex and is fitting the noise in the data instead of the underlying patterns. This can result in poor performance on new, unseen data.